

# Introduction

to

# Virtualization

## NREN

Feb 10 - 12, 2023



This material is licensed under a Creative Commons Attribution-NonCommercial 4.0 International License (<http://creativecommons.org/licenses/by-nc/4.0/>)

# What is Virtualization?

- Virtualization is the creation of a virtual version of something, such as
  - a operating system
  - a server
  - storage devices
  - network resources
- It allows multiple systems to run on a single physical machine, increasing resource utilization and improving hardware independence.

# Terminology

- **Virtualization:**
  - dividing available resources into smaller independent units
- **Emulation:**
  - using software to simulate hardware which you do not have
- The two often come hand-in-hand
  - e.g. we can *virtualize* a PC by using it to *emulate* a collection of less-powerful PCs

# Virtualization: Benefits

- **Consolidation**

- Most systems are under-utilized, especially the CPU is idle for much of the time
- Do more work with less hardware
- Reduced space and power requirements

- **Management**

- Less hardware inventory to manage
- Concentrate your resilience efforts
- Increased isolation between services
- Abstract away (hide) differences in hardware

# Virtualization: Benefits (contd.)


- **Flexibility**

- Grow systems on demand (e.g. allocate more CPU or RAM where it is needed)
- Create new services quickly without having to install new hardware every time
- Dynamically create and destroy instances for testing and development

- **New capabilities**

- Snapshot/restore, cloning, migration, ...
- Run different OSes on the same machine at once

# Virtualization: a familiar example




- Who has not seen this before?!
- Like having two (or more) hard drives
  - you get to choose the sizes
- **Why** is this useful?

# Virtualization: Another example

- Virtualize a switch: VLANs
  - like dividing a switch into separate switches
- Benefits:
  - can keep traffic separate (broadcast domains)
  - can create VLANs and how they are assigned to ports, purely through software configuration
  - can combine VLANs onto a single cable and split them out again (tagging/trunking)

# VLANs




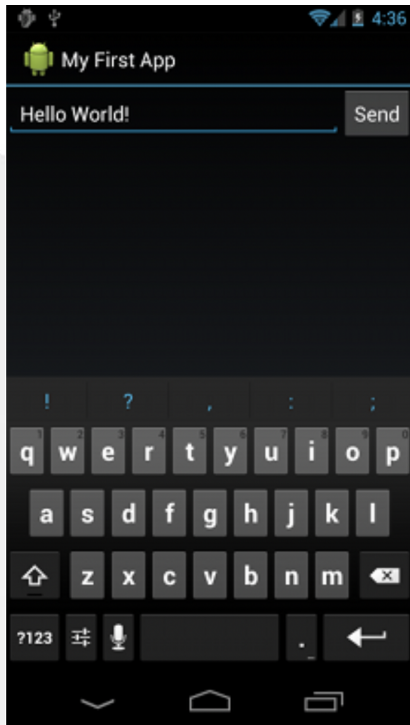


# Emulation

- In software, you can simulate the behaviour of a device which doesn't exist
- Example: emulation of a CD-ROM drive using an ISO file
  - a request to read block N of the (virtual) CD-ROM drive instead reads block N of the ISO file
  - similar to partition mapping
- You can simulate any hardware - including the CPU or an entire system!

# Emulation: example

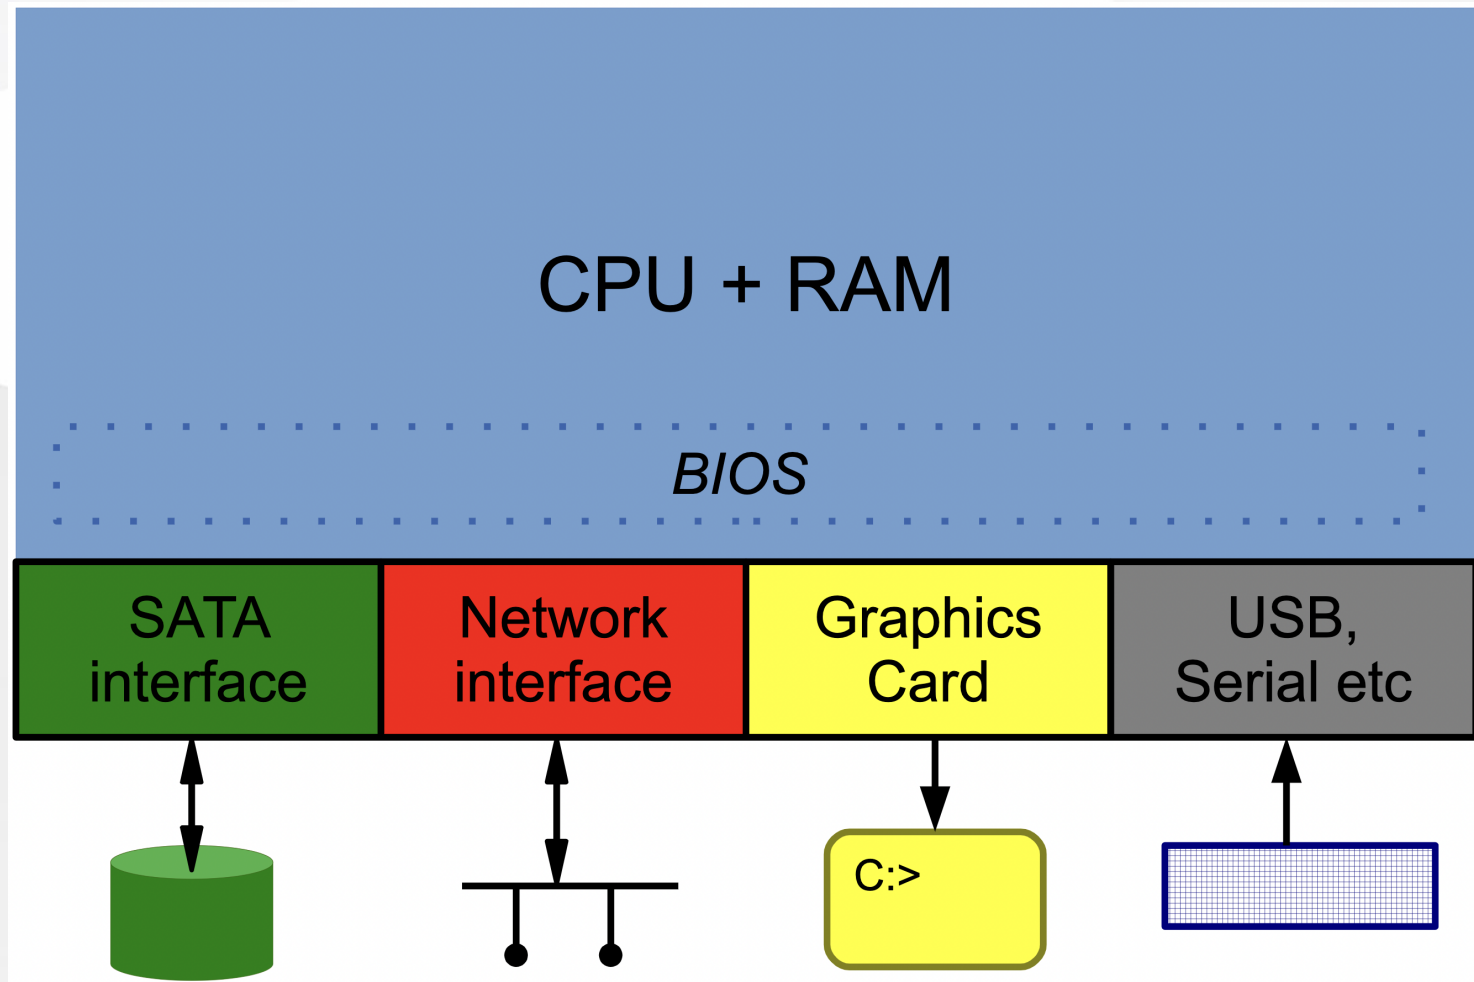
- Android SDK
  - Emulates an Android smartphone with ARM CPU
  - The "screen" is mapped to a window on your PC



# Emulation: more examples

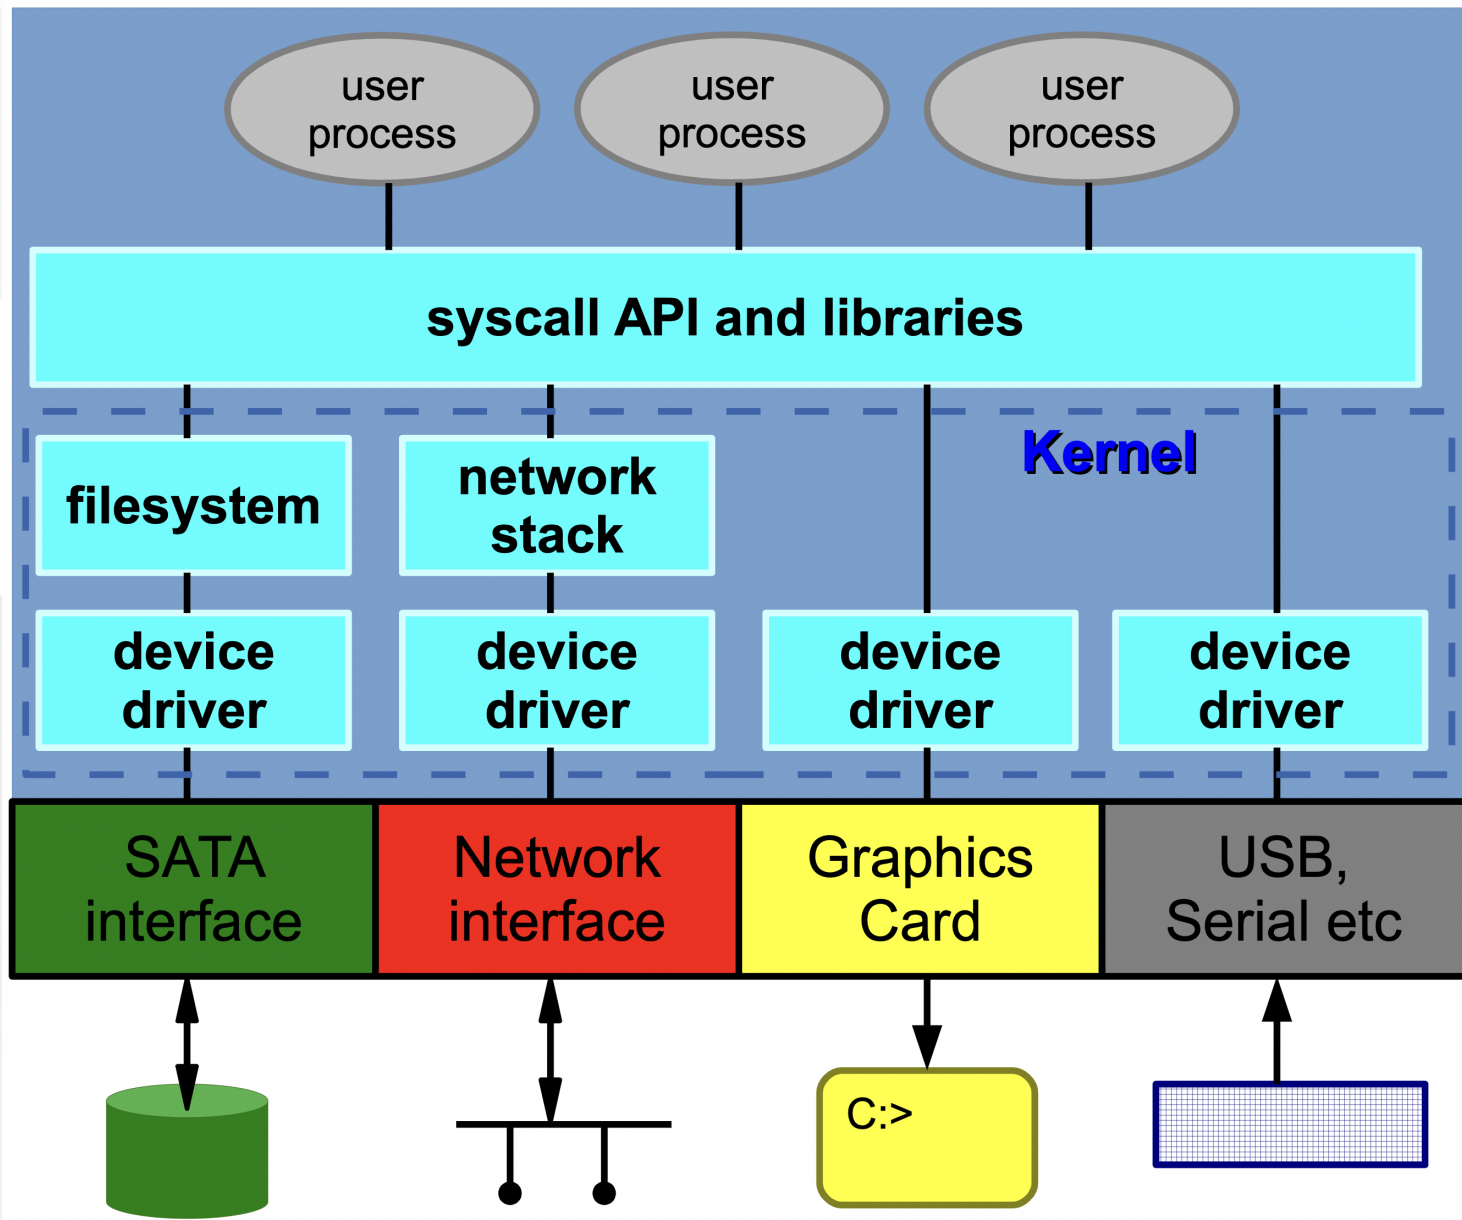
- Dynamips / Dynagen / GNS3
  - Emulates a Cisco router with MIPS CPU and network interfaces
- QEMU
  - Emulates an entire PC (i386 processor and interfaces)

# What's in a PC?



# Boot up sequence

- A small program (the BIOS) runs when machine is switched on
- It uses the hardware to load an operating system—boot from hard drive, USB/CD-ROM, network...
- Modern operating systems then ignore the BIOS from that point onwards
- The next slide shows a machine after it has booted up (simplified)



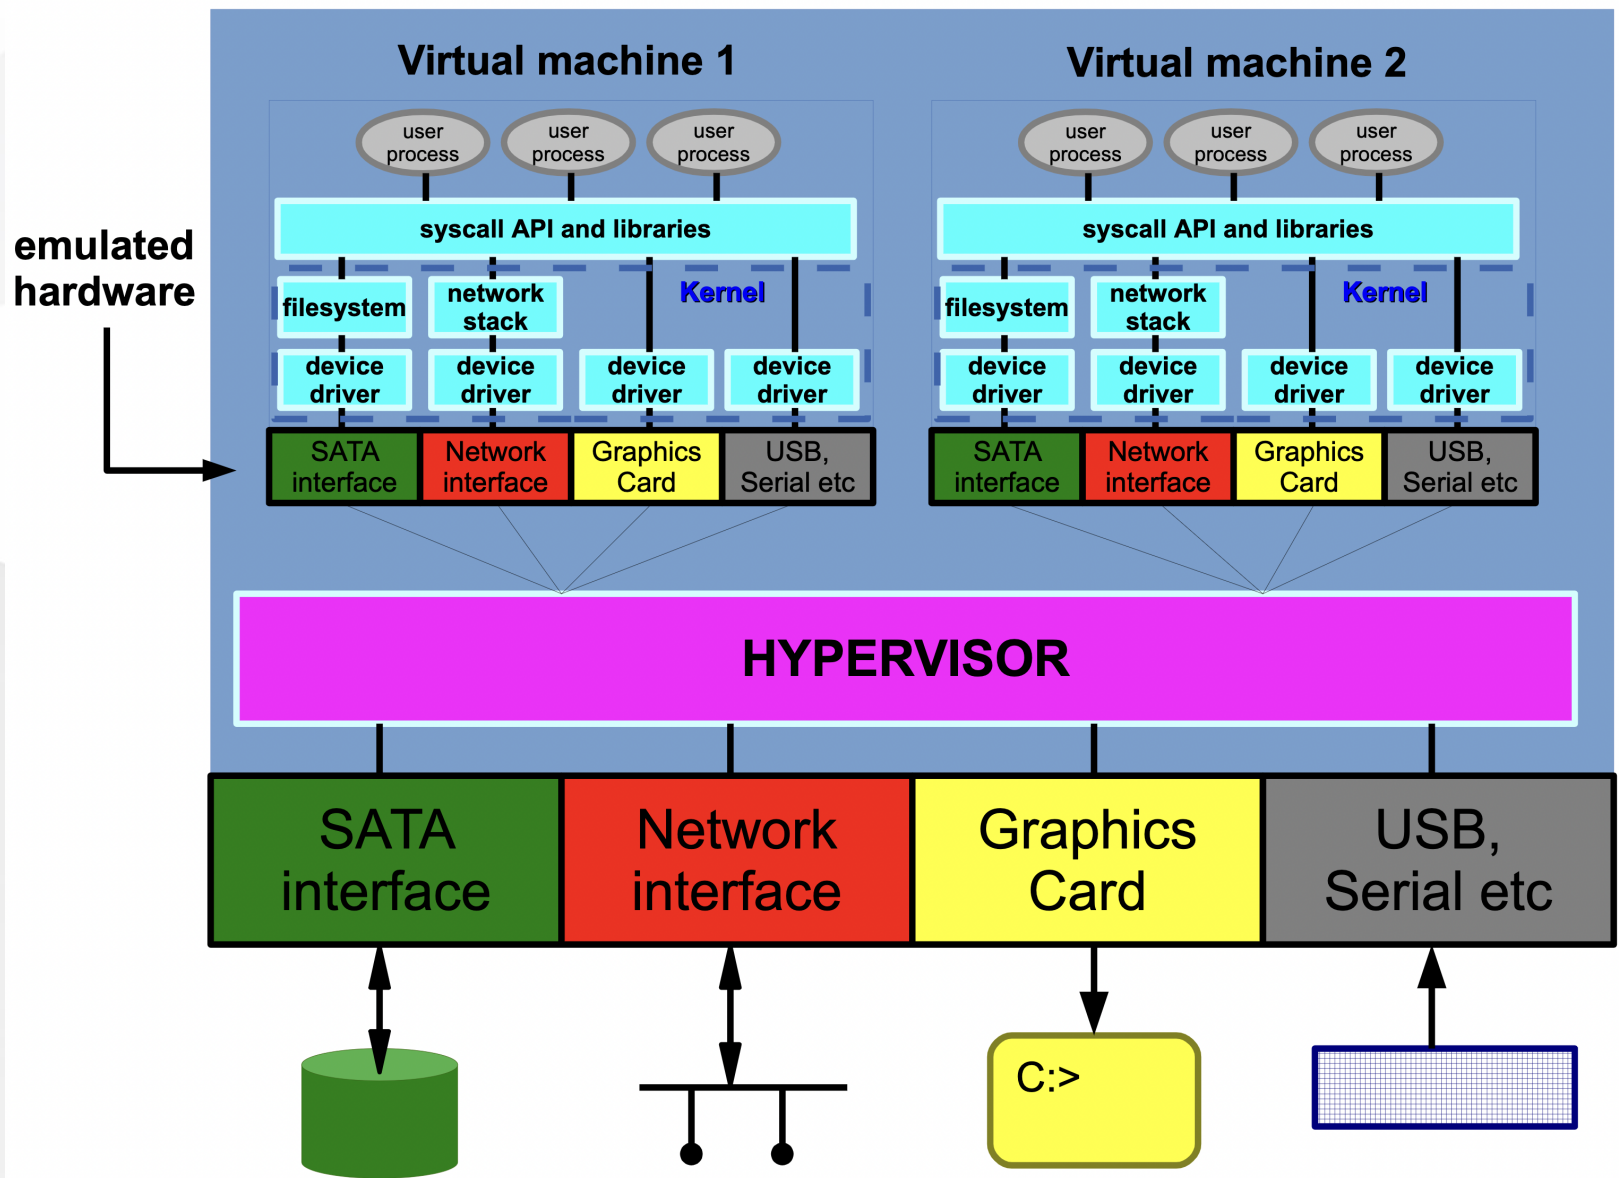
# Points to note

- The device drivers in the OS interact with the hardware
- User processes are forbidden by the OS from interacting directly with the hardware
  - the OS configures protection mechanisms to enforce this

# What we need to emulate/virtualize a PC?

- we must emulate all the components of the PC
  - hard disk interface, network card
  - graphics card, keyboard, mouse
  - clock, memory management unit etc
- We want multiple instances to co-exist and not be able to interfere with each other
  - access to memory must also be controlled
- The software to do this is called a ***hypervisor***






# Virtual Machines

- Each emulated PC is a "**virtual machine**"
- Hypervisor allocates some real system RAM to each VM, and shares the CPU time
- Hypervisor emulates other hardware, e.g. disk and network interfaces
- Within each VM you can boot an operating system
- Full hardware virtualization means different VMs can be running different OSes

# Virtualization terminology

- The **host** is the machine running the emulation software
- The **guest** is the emulated (virtual) machine
- One host could be running many guests



# The Hypervisor

- a software layer that sits between the operating system and physical hardware of a computer
- its purpose is to create virtual machines that can run multiple, isolated operating systems on a single physical machine.
- it manages and allocates the underlying physical resources to each virtual machine.
- serves as the intermediary between the virtual machines and the physical resources, ensuring that each virtual machine has access to the resources it needs to operate

# Types of Hypervisor

- Type 1 hypervisors
- Type 2 hypervisors

# Type 1 hypervisors

- known as bare-metal hypervisors
- run directly on the host's physical hardware
- provide a virtualized environment for guest operating systems
- examples:
  - VMware ESXi
  - Microsoft Hyper-V
  - Proxmox Virtual Environment
  - XCP-ng
  - Xen

# Type 2 hypervisors

- known as hosted hypervisors
- run on a host operating system
- provide a virtual environment for guest operating systems
- examples:
  - Oracle VirtualBox
  - VMware Workstation
  - VMware Fusion
  - Parallels Desktop

# Emulated disk hardware

- A hard drive is a "block device"
  - OS makes requests like "read block number 42", "write block number 99"
- Real hard drives have a fixed size!
  - This is what the guest OS will expect to see
- So the hypervisor must redirect these accesses to something else



# Emulated disk hardware (contd.)

- Options include:
  - a disk image file on the host (simple)
  - a partition or logical volume on the host (faster)
  - a remote file or remote block device (via network)
- A disk image file is easy to backup and transfer from host to host
- There are different ways to make a disk image file. Suppose we want the guest to see a 10GB virtual hard drive?
  - raw
  - image file

# RAW File or Disk

- A "raw" file is just a plain 10GB data file
  - Nth block of the virtual hard drive corresponds to the Nth block in the image file
  - if this is allocated up-front, you use 10GB of (hopefully) contiguous space on the host
  - Fast in operation, avoids fragmentation on the host
  - Wasteful of space
  - Slow to create
  - Slow to copy

# Image File

- Custom VM image format with header and data
  - can be Thin provisioned
    - doesn't allocate space until each block is written to
    - reading from unallocated space reads zeros
    - can leads to fragmentation
    - can lead to failures if filesystem becomes full
  - can be Thick provisioned
    - pre-allocate all the space
    - wasteful of space
    - slow to create
- Various formats, e.g. VDI (virtualbox), VMDK (VMware), QCOW2 (qemu/kvm)
- Also add features like snapshots

# Emulated network hardware

- Each guest NIC gets a fake MAC address
- Different ways to interconnect with host NIC
- NAT
  - outbound packets translated to share the host's IP address
- Bridging
  - packets sent/received untranslated over the host's physical NIC
  - Each VM gets its own IP address on the external network
  - More transparent
  - Does not always work on wireless NICs though

# Summary

- Virtualization can make better use of your hardware by emulating more machines than you really have
- The emulated environment is provided by a ***hypervisor***
- The hypervisor (host) lets you start up virtual machines (guests) each with its own operating system and emulated devices
- Guest hardware emulated using resources on the host

